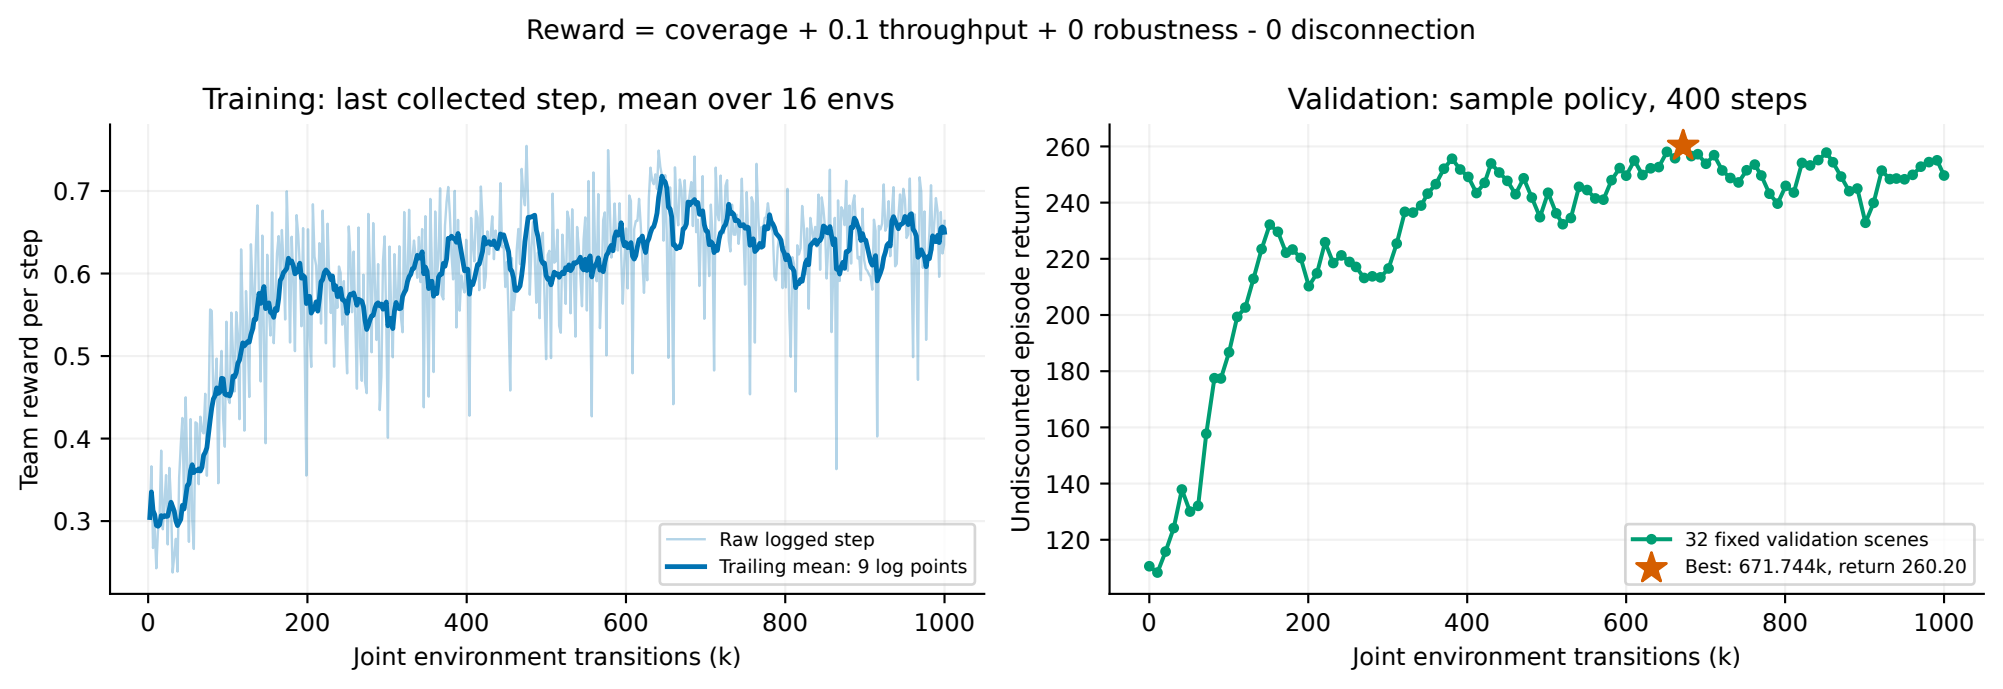
Reward = coverage + 0.1 throughput + 0 robustness - 0 disconnection
Training: last collected step, mean over 16 envs
Team reward per step
0.3
0.4
0.5
0.6
0.7
0
200
400
600
800
1000
Joint environment transitions (k)
Raw logged step
Trailing mean: 9 log points
Validation: sample policy, 400 steps
Undiscounted episode return
120
140
160
180
200
220
240
260
0
200
400
600
800
1000
Joint environment transitions (k)
32 fixed validation scenes
Best: 671.744k, return 260.20